

# Probability

Jacopo Notarstefano

`jacopo.notarstefano [at] gmail.com`

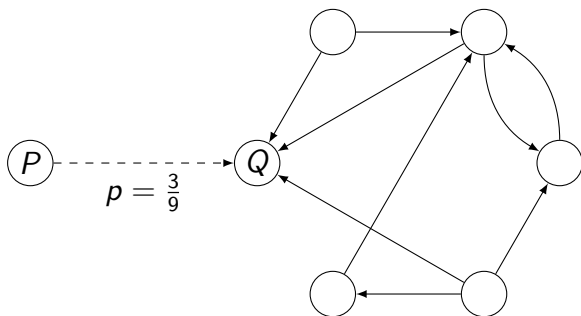
11 February 2014

# Main ideas

- ➊ Given two pages, we want to estimate how many pages would be linked by both if links were created randomly.
- ➋ If the actual number is smaller, then we conclude that those pages are not related. If it's bigger, we assign a score between 0 and 1.
- ➌ This estimate depends on how we model random link creation between pages.

# The Barabási–Albert model

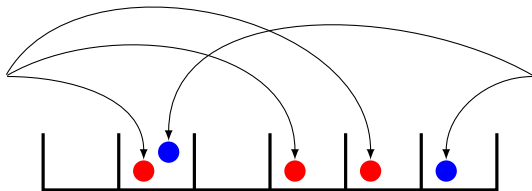
The probability that a new page  $P$  links an existing page  $Q$  is proportional to  $\text{indeg}(Q)$ : “The rich get richer”.



# Balls and bins, 1/2

## Problem

*Suppose that we have  $W$  bins,  $n_1$  red balls and  $n_2$  blue balls. When we throw a ball it falls in bin  $i$  with probability  $p_i$ . When we are done throwing all the balls, what's the expected number of bins with both a blue and a red ball?*



# Balls and bins, 2/2

## Solution

*If **all throws are independent**, then, by linearity of expectation, we have*

$$\mathbb{E}[|N_1 \cap N_2|] = \sum_{i,j=1}^{n_1, n_2} \mathbb{E}[I_{ij}] = n_1 n_2 \sum_{i=1}^W p_i^2 = n_1 n_2 \mathbf{P}$$

*where  $I_{ij}$  is random indicator variable denoting that red ball  $i$  and blue ball  $j$  landed in the same bin.*

# The algorithm

---

## Algorithm 1 Probability

---

```
// Preprocessing step
Scan Wikipedia and compute P
// For each pair of pages  $P_1$  and  $P_2$ 
 $N_1 \leftarrow \text{outLinks}(P_1); n_1 \leftarrow N_1.\text{length}$ 
 $N_2 \leftarrow \text{outLinks}(P_2); n_2 \leftarrow N_2.\text{length}$ 
 $\text{actualValue} \leftarrow |N_1 \cap N_2|$ 
 $\text{expectedValue} \leftarrow n_1 n_2 \cdot \mathbf{P}$ 
if  $\text{actualValue} < \text{expectedValue}$  then
    return 0
else
    return  $\text{normalize}(\text{actualValue} - \text{expectedValue})$ 
end if
```

---

# Results

The algorithm manages to retain the recall baseline while improving on its precision, thus achieving a better F1 score.

	TagMe baseline	0.562	0.586	0.540
	group3	0.594	0.660	0.539

The algorithm is also fast: a run against the entire AIDA/CoNLL dataset takes less than 2 minutes.